

Assignment 5.2

October 3, 2022

```
[1]: from keras.datasets import reuters
      (train_data, train_labels), (test_data, test_labels) = reuters.load_data(
      num_words=10000)
```

```
[2]: len(train_data)
```

```
[2]: 8982
```

```
[3]: len(test_data)
```

```
[3]: 2246
```

```
[4]: train_data[10]
```

```
[4]: [1,
      245,
      273,
      207,
      156,
      53,
      74,
      160,
      26,
      14,
      46,
      296,
      26,
      39,
      74,
      2979,
      3554,
      14,
      46,
      4689,
      4329,
      86,
      61,
```

```
3499,  
4795,  
14,  
61,  
451,  
4329,  
17,  
12]
```

```
[5]: train_labels[10]
```

```
[5]: 3
```

```
[6]: import numpy as np  
  
def vectorize_sequences(sequences, dimension=10000):  
    results = np.zeros((len(sequences), dimension))  
    for i, sequence in enumerate(sequences):  
        results[i, sequence] = 1.  
    return results  
x_train = vectorize_sequences(train_data)  
x_test = vectorize_sequences(test_data)
```

```
[7]: # one hot encoding  
  
def to_one_hot(labels, dimension=46):  
    results = np.zeros((len(labels), dimension))  
    for i, label in enumerate(labels):  
        results[i, label] = 1.  
    return results  
one_hot_train_labels = to_one_hot(train_labels)  
one_hot_test_labels = to_one_hot(test_labels)
```

```
[8]: # Define the model  
  
from keras import models  
from keras import layers  
model = models.Sequential()  
model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))  
model.add(layers.Dense(64, activation='relu'))  
model.add(layers.Dense(46, activation='softmax'))
```

```
[9]: # Compiling the model  
  
model.  
    ↪ compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
[10]: # Validation set
```

```
x_val = x_train[:1000]
partial_x_train = x_train[1000:]
y_val = one_hot_train_labels[:1000]
partial_y_train = one_hot_train_labels[1000:]
```

```
[11]: # Training the model
```

```
history = model.
    ↳fit(partial_x_train,partial_y_train,epochs=20,batch_size=512,validation_data=(x_val,
    ↳y_val))
```

Epoch 1/20

16/16 [=====] - 1s 31ms/step - loss: 2.6418 - accuracy: 0.5396 - val_loss: 1.7667 - val_accuracy: 0.6720

Epoch 2/20

16/16 [=====] - 0s 21ms/step - loss: 1.4323 - accuracy: 0.7076 - val_loss: 1.3215 - val_accuracy: 0.7130

Epoch 3/20

16/16 [=====] - 0s 19ms/step - loss: 1.0533 - accuracy: 0.7803 - val_loss: 1.1295 - val_accuracy: 0.7610

Epoch 4/20

16/16 [=====] - 0s 18ms/step - loss: 0.8199 - accuracy: 0.8295 - val_loss: 1.0222 - val_accuracy: 0.7920

Epoch 5/20

16/16 [=====] - 0s 17ms/step - loss: 0.6480 - accuracy: 0.8657 - val_loss: 0.9836 - val_accuracy: 0.7920

Epoch 6/20

16/16 [=====] - 0s 18ms/step - loss: 0.5155 - accuracy: 0.8943 - val_loss: 0.9169 - val_accuracy: 0.8060

Epoch 7/20

16/16 [=====] - 0s 15ms/step - loss: 0.4120 - accuracy: 0.9148 - val_loss: 0.8927 - val_accuracy: 0.8150

Epoch 8/20

16/16 [=====] - 0s 17ms/step - loss: 0.3357 - accuracy: 0.9305 - val_loss: 0.8979 - val_accuracy: 0.8170

Epoch 9/20

16/16 [=====] - 0s 17ms/step - loss: 0.2772 - accuracy: 0.9389 - val_loss: 0.8985 - val_accuracy: 0.8220

Epoch 10/20

16/16 [=====] - 0s 21ms/step - loss: 0.2334 - accuracy: 0.9445 - val_loss: 0.8989 - val_accuracy: 0.8200

Epoch 11/20

16/16 [=====] - 0s 18ms/step - loss: 0.2018 - accuracy: 0.9509 - val_loss: 0.9652 - val_accuracy: 0.8080

Epoch 12/20

```

16/16 [=====] - 0s 18ms/step - loss: 0.1822 - accuracy:
0.9524 - val_loss: 0.9262 - val_accuracy: 0.8290
Epoch 13/20
16/16 [=====] - 0s 20ms/step - loss: 0.1599 - accuracy:
0.9559 - val_loss: 1.0184 - val_accuracy: 0.8040
Epoch 14/20
16/16 [=====] - 0s 16ms/step - loss: 0.1485 - accuracy:
0.9574 - val_loss: 0.9404 - val_accuracy: 0.8240
Epoch 15/20
16/16 [=====] - 0s 16ms/step - loss: 0.1382 - accuracy:
0.9569 - val_loss: 1.0015 - val_accuracy: 0.8080
Epoch 16/20
16/16 [=====] - 0s 15ms/step - loss: 0.1251 - accuracy:
0.9584 - val_loss: 0.9891 - val_accuracy: 0.8180
Epoch 17/20
16/16 [=====] - 0s 18ms/step - loss: 0.1249 - accuracy:
0.9572 - val_loss: 1.0377 - val_accuracy: 0.8110
Epoch 18/20
16/16 [=====] - 0s 16ms/step - loss: 0.1209 - accuracy:
0.9560 - val_loss: 1.0296 - val_accuracy: 0.8170
Epoch 19/20
16/16 [=====] - 0s 16ms/step - loss: 0.1121 - accuracy:
0.9590 - val_loss: 1.0997 - val_accuracy: 0.8020
Epoch 20/20
16/16 [=====] - 0s 16ms/step - loss: 0.1131 - accuracy:
0.9592 - val_loss: 1.0645 - val_accuracy: 0.8080

```


```

[12]: # Plot training and validation loss

import matplotlib.pyplot as plt


loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(loss) + 1)
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()

```



```
[13]: # Plotting the training and validation accuracy

plt.clf()
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
[14]: # Retraining a model from scratch

model = models.Sequential()
model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(46, activation='softmax'))
model.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(partial_x_train, partial_y_train, epochs=9, batch_size=512, validation_data=(x_val,
→ y_val))
results = model.evaluate(x_test, one_hot_test_labels)
```

```
Epoch 1/9
16/16 [=====] - 1s 33ms/step - loss: 2.5460 - accuracy:
0.5307 - val_loss: 1.6932 - val_accuracy: 0.6450
Epoch 2/9
16/16 [=====] - 0s 21ms/step - loss: 1.4022 - accuracy:
0.7068 - val_loss: 1.3101 - val_accuracy: 0.7140
Epoch 3/9
16/16 [=====] - 0s 20ms/step - loss: 1.0643 - accuracy:
0.7714 - val_loss: 1.1514 - val_accuracy: 0.7520
Epoch 4/9
16/16 [=====] - 0s 19ms/step - loss: 0.8361 - accuracy:
```

```
0.8235 - val_loss: 1.0282 - val_accuracy: 0.7810
Epoch 5/9
16/16 [=====] - 0s 18ms/step - loss: 0.6711 - accuracy:
0.8592 - val_loss: 0.9590 - val_accuracy: 0.7990
Epoch 6/9
16/16 [=====] - 0s 16ms/step - loss: 0.5377 - accuracy:
0.8866 - val_loss: 0.9405 - val_accuracy: 0.7900
Epoch 7/9
16/16 [=====] - 0s 16ms/step - loss: 0.4303 - accuracy:
0.9104 - val_loss: 0.9067 - val_accuracy: 0.8160
Epoch 8/9
16/16 [=====] - 0s 20ms/step - loss: 0.3544 - accuracy:
0.9257 - val_loss: 0.9113 - val_accuracy: 0.8080
Epoch 9/9
16/16 [=====] - 0s 25ms/step - loss: 0.2933 - accuracy:
0.9342 - val_loss: 0.9195 - val_accuracy: 0.8160
71/71 [=====] - 0s 3ms/step - loss: 0.9927 - accuracy:
0.7801
```

```
[15]: results
```

```
[15]: [0.9926513433456421, 0.780053436756134]
```

```
[16]: import copy
test_labels_copy = copy.copy(test_labels)
np.random.shuffle(test_labels_copy)
hits_array = np.array(test_labels) == np.array(test_labels_copy)
float(np.sum(hits_array)) / len(test_labels)
```

```
[16]: 0.18432769367764915
```

```
[17]: # Generating predictions for new data

predictions = model.predict(x_test)
predictions[0].shape
```

```
[17]: (46,)
```

```
[18]: np.sum(predictions[0])
```

```
[18]: 0.9999999
```

```
[19]: np.argmax(predictions[0])
```

```
[19]: 3
```